

# CSci 4061

## Introduction to Operating Systems

### Processes in C/Unix

#### Chapter 3 (R&R)

# Process as Abstraction


- Talked about C programs a bit
- Program is a static entity
- Process is an abstraction of a running program provided by the OS
  - Granted resources

# Process Abstraction

- Process operations
  - Create: `fork`
  - Change: `exec*`
  - Terminate: `exit/abort/return/signals`
  - Synchronize: `wait*`, and others (later)

# Virtual Address Space

- Memory map of the process
- Virtual address space (VAS)
  - Set of legal addresses

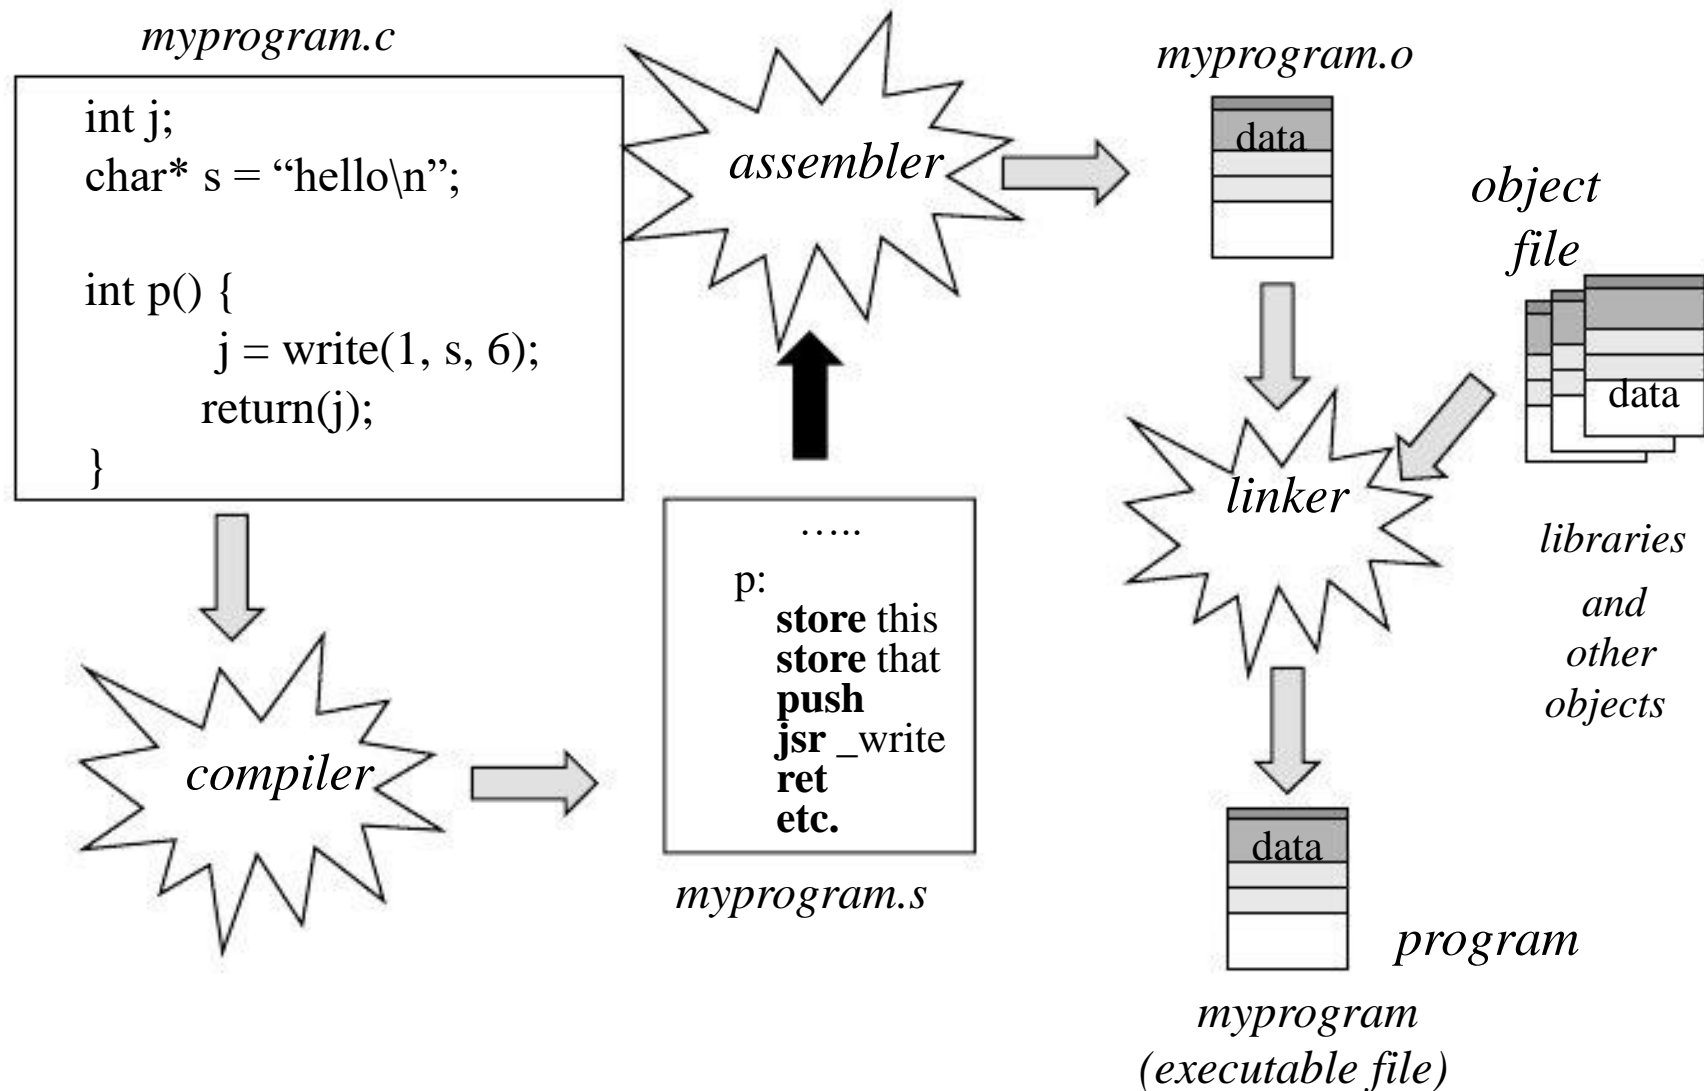


Address space disjoint

Isolation

# Top-down: Why Processes?

# The Birth of a Process: Executable




# Unix Processes

- Process trees
- Creation: `fork/exec`
- Synchronization/Control: `wait`
- Termination: `exit`
- Error handling
- Identities

# Processes in Unix

- Processes in Unix form a hierarchy
  - Root is called *systemd*, started by OS
  - Child is created by parent
  - Relationship important for communication
- Unix shell is a child of *systemd*, shell in turn creates processes





# “Your process”

- When you log in ...
- A process known as your shell is created and running
- Has an associated terminal window for I/O
- It is a foreground process and allows you to interact with it

# Processes in Unix (cont'd)

```
shell> cat file1 file2
```

... output

```
shell>
```

Example of a **foreground** process: shell waits for completion + user can interact with process

Under the hood:

- shell creates “cat” process
- waits for it to finish
- prompts for next command

# Processes in Unix (cont'd)

```
shell>my_program& // & put in the  
background
```

```
shell>cat foo.c
```

**Background process:** shell does not wait for `my_program` to complete before next prompt; user cannot interact with background process

Examples?

# Switching: foreground to background

- `<loop>`
- `^z` to suspend a process; get control back
- `bg` run it in the background
- `fg` run it in the foreground
- process shell creates -- *job*
- `jobs`: lists job ids (different from process ids)  
in current shell only (`top`, `ps` machine-wide)

# System Programming Interface: `fork`

- The Unix system call for process creation is called `fork()`

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
pid_t fork(); // pid_t is a process id (#)
```

- With any Unix system call: error codes, header files, parameters
- The `fork` system call creates a “clone” of the parent process

# Fork (cont'd)

- Child is given a “copy” of the parent’s memory (virtual address space)
  - actually a virtual copy using copy-on-write
- Child is running the same program code as the parent
- Child begins life with the same register values of the parent (e.g. PC)
- Child inherits resources from the parent
  - open files
- What is not shared?
  - locks, IDs, signals, CPU time measures, ...

## PARENT

```
#include <unistd.h>
int pid;
int status = 0;
pid = fork ();
// Parent: PID IS NON-ZERO
if (pid > 0) {
    printf ("Parent: child has
            pid=%d", pid);

    ....
    pid = waitpid(pid, &status, 0);
} else if (pid == 0) {
    printf ("child here");
    ....
    exit(status);
} else {
    perror ("fork problem");
    exit (-1);
}
```

Entry point: main

## CHILD


```
#include <unistd.h>
int pid;
int status = 0;
pid = fork ();
// HERE IS WHERE WE START, pid = 0
if (pid > 0) {
    printf ("Parent: child has
            pid=%d", pid);

    ....
    pid = waitpid(pid, &status, 0);
} else if (pid == 0) {
    printf ("child here");
    ....
    exit(status);
} else {
    perror ("fork problem");
    exit (-1);
}
```

Entry point: after fork


# Running programs: memory and the CPU

CPU



$x == y$

address of fork  $z$



$x$

Program B

$x$ : fork

Data



address

main memory



# Example: Process Creation via `Fork`

```
#include <unistd.h>
int pid;
int status = 0;
pid = fork ();
if (pid > 0) {
    printf ("Parent: child has
            pid=%d", pid);

    ...
    pid = waitpid(pid, &status, 0);
} else if (pid == 0) {
    printf ("child here");
    ...
    exit(status);
} else {
    perror ("fork problem");
    exit (-1);
}
```

`Fork` returns *twice*. It returns a 0 PID to the child and the child's PID to the parent

Parent typically **blocks** or *waits* until the child terminates by using `wait` or `waitpid`

Child or any process can return an exit status

Always check for errors (-1) on syscalls

# Fork example

simplefork.c:

- `while (1) fork ();`
  - fork “bomb”

# Process Topologies

```
#include <sys/types.h>
#include <unistd.h>
pid_t fork(); // pid_t is a process id (#)
```

- simplefork.c
- simplechain.c

# Waiting

Called by parent to `wait` (**block**) until child exits

```
#include <sys/wait.h>
```

```
pid_t wait (int *stat_loc); // any child
```

```
pid_t waitpid (pid_t pid, int *stat_loc,  
               int options); // spec. child
```

`stat_loc`: {why} did child exit? (don't care, NULL or 0)

return parameter, if not NULL, must allocate it!

```
wait (NULL); // OR
```

```
int stat_loc;
```

```
wait (&stat_loc); // "output" param
```

What is in `stat_loc`? See 3.4.1 Status values (R&R)

1. status of child: exited, suspended, ...

2. If normal exit, what was the exit value (i.e. child does `return (5);`);

# Processes and shared variables

- Processes only share at CLONE time - copy

```
int i = 5;
```

```
childpid = fork();
```

```
if (childpid == 0) {
```

```
    print i; // #1
```

```
    i = 7;
```

```
    print i; // #2
```

```
}
```

```
else {
```

```
    print i;    // #3
```

```
    i = 3;
```

```
    print i;    // #4
```

```
    wait (NULL);
```

```
    print i;    // #5
```

```
}
```

```
/* child code */
```

```
/* parent code */
```

# Top-down: Why Processes?

- Why we need fork? Multi-tab browser.

```
// #define MaxURL 100
const int MaxURL = 100;
int i, num_urls;
char *URLs [MaxURL];
pid_t child_pids [MaxURL];
// assume num_urls and URLs have been set
for (i=0; i<num_urls; i++)
    if ((child_pids[i] = fork()) == 0) {
        fetch_and_display (URLs[i]);
        break;
    }
for (i=0; i<num_urls; i++) {
    wait (NULL);
    // why wait? cleanup any state related to tabs
}
```

**ISOLATION!**  
for tabs

# How can Fork fail?

Too many processes in the system

Not enough memory

Not enough disk space (later)

# Back to `Fork` (cont'd)

- The child process may execute a different *program*: `exec*` call
- Still need `fork` to create the child clone and container, but `exec*` changes the code it runs
- It completely over-writes the child's memory
  - Except: IDs (`pid`) are intact as are I/O descriptors, environment



# Why do we need exec?

// implementation of shell

```
> ./foo
```

```
> gcc ...
```

// implementation of make

```
foo.o: foo.h foo.c
```

```
gcc foo.c
```

```
cp foo.o /usr/bin/foo.o
```

```
...
```

# execl

```
int execl (const char *path,  
          const char *arg0,  // strings  
          const char *arg1, ...  
          (char *) 0);
```

Executes program named by 1<sup>st</sup> argument

Pathname+executable name (e.g. “/usr/jon/prog”)

arg0 is just the name of the executable “prog”

remainder are optional, with 0 or NULL terminating

# Many Options

- Family of options
  - `execl`, `execvp`, `execv`, `execvp`, `execve`
  - Differ in how arguments + env are used
- All short forms of `execve`

```
int execve (const char *path,  
            const char *argv[],  
            const char *envp[]);
```

provides new values for env



# Example:

<execcmd.c>

How can `exec*` fail?

# Why do we need `exec*`:

## Consider the Shell

```
while (TRUE) {                                /* repeat forever */
    type_prompt( );                            /* display prompt */
    read_command (command, parameters)        /* input from terminal */
    /* some error checking ... */
    if (fork() != 0) {                        /* fork off child process */
        /* Parent code */
        waitpid( -1, &status, 0);            /* wait for any child to exit */
    } else {
        /* Child code */
        execve (command, parameters, 0);      /* execute command */
    }
}
```

If command crashes, no problem!

# Shell with `execve`

- When you run your program:

```
shell> my_prog 0 10
```

The Shell is doing the `fork/exec`

path/command is *current working directory* + `my_prog`

`argv:`

`arg0` is `"my_prog"`

`arg1` is `"0"`

`arg2` is `"10"`

# Questions

```
read_command (command, parameters)
if (fork() != 0) {
    /* Parent code */
    waitpid( -1, &status, 0);
} else {
    execve (command, parameters, 0);
}
}
```

How would you implement:

1) `shell> my_prog&`


2) what would happen if the shell did not fork, just exec?

# Uses of exec\*

- Program that runs other programs



# Putting it all together



# Process Termination


- How does a process terminate?
  - Return from `main`
  - Falls off the end of `main`
  - Call `exit`
  - Call `abort`
  - Receives a death signal or exception
    - `kill (pid_t pid, SIGKILL)`

# Process Termination (cont'd)

- Should be orderly
  - Child terminates before parent and parent is `waiting`
- If child `exits` when parent is not `waiting`
  - Child becomes a “zombie”
  - OS keeps it around long enough until parent does a `wait` to get `exit` status or until parent `exits`

# Process Termination (cont'd)

- Parent exits while child is running
  - Child becomes an orphan
  - No problem as children become “adopted” by parent up the tree, possibly *systemd*



# Identities

- When a process runs, OS must associate a user and group id to it, why?
  - accounting and security

`fork ()` returns pid of the child to the parent

`getpid ()` returns the pid of the calling process

`getppid ()` returns my parents pid

`getuid ()` returns userid of the user that started the process (e.g. “jon” => 89392)

# Shell and &

- Does the shell ever wait for background processes?
- Yes, it receives a signal (SIGCHLD) when child exists (to be discussed later)

# Next Time

- Begin discussion of input/output
- Read Chapter 4 R&R
- Project #1 will be available on Weds