

CSci 4061

Introduction to Operating Systems

Module 2: Input/Output

Input/Output: High-level

I/O Topics

- First, cover high-level I/O
- Next, talk about low-level device I/O
- I/O not part of the C language!

I/O Device Abstraction(s)

- Abstraction: source/sink for data (raw bytes)
- Operations:
 - Open/close device
 - Read from device
 - Write to device
 - Control device

High-level I/O

- Go further with the abstraction
- Hide “device” or low-level I/O
- Low-level I/O abstraction: source/sink for data
- High-level: more than raw bytes --
- Features
 - Stream abstraction
 - Formatted/typed I/O
 - String-based I/O
 - Line-oriented buffering

Streams and Files

- Unix `FILE` object
 - Delivers an ordered sequence of bytes
 - Defined in `<stdio.h>`
 - User-space library
 - Built on top of low-level file descriptors
- Three default streams
 - `stdin`
 - `stdout`
 - `stderr`

Inside a FILE

- Points to actual file
- Current offset
- Mode : read, write, append, etc.
- Buffers
- Write Buffering (open for write)
 - Internal character buffer of size BUFSIZE
 - Writes are done to in-memory buffer
 - When buffer is full (BUFSIZE), `write` out buffer
 - Or if line-buffered stream (e.g. `stdout`) when `'\n'` is written

Inside a FILE (cont'd)

- Read Buffering (open for read)
 - Reads are done from in-memory read buffer
 - When buffer is empty, we `read` a chunk of BUFSIZE

--	--	--	--
- Why is the use of an internal buffer for I/O advantageous?

Buffering: FFLUSH

- Sometime we want to “force” buffer to spill to the OS - why?
 - see data immediately (OS auto-flushes K sec)

```
int fflush (FILE *stream);
```

```
fprintf (F, "your output is %d");
```

```
fflush (F); // force it
```

`stderr` always causes buffer to be flushed

`stdout` with ‘\n’ as well

OPEN/CLOSE

```
#include <stdio.h>
```

```
// NULL on failure, errno also set
```

“w” will create file
if not there ...

```
FILE *fopen (const char *filename,  
             const char *mode);
```

```
mode = “r”, “w”, “a” (others, r+, w+)
```

 truncate to 0 length

```
// returns 0 on success, EOF otherwise (-1)
```

```
int fclose (FILE *stream);
```

I/O Access

- Assumed to be sequential
- The `FILE*` keeps track of the current file offset for read and write

Character-based I/O

```
#include <stdio.h>
// SUBSET of calls: return EOF on failure or end
int fgetc (FILE *stream);
int fputc (int char, FILE *stream);
int getc (); // STDIN
int putc (int char); // STDOUT
```

Which header are things located?

>man getc

String-based I/O

// reads `nsiz-1` characters or up to a newline

// if newline, then `'\n'` goes into `buf`

// appends `'\0'` at the end of `buf`

// **caller allocs** `buf`

// returns `buf` or `NULL` if at EOF or an error occurs

`char *fgets (char *buf, int nsiz, FILE *inf);`

sets `errno`



// outputs `buf` - better be `'\0'` terminated **what must be true of `buf`?**

// returns last char written (+ `#`) or EOF if an error occurs

`int fputs (const char *buf, FILE *outf);`

sets `errno`



Also: `gets`, `puts`

Importance of newline ‘\n’

- `stdin` is a line-oriented device
 - ‘\n’ (EOL), input is read when ‘\n’ is seen or EOF
 - ^D is EOF
- `fgets` on `stdin` - will include ‘\n’ in the string!
- `gets` does not

String library functions

- Not system calls
- On your own: remember to allocate, ‘\0’
- `strcpy`, `strncpy`, `strlen`, `strtok`,
`strcmp`, ..., *strdup*

What about I/O for data types?

- Formatted I/O!
- Formatted output and “output” strings

```
int fprintf (FILE *outf,  
            const char *fmt,  
            args)
```

```
int sprintf (char *string, // alloc!  
            const char *fmt,  
            args)
```

Also: `printf`

Formatted Output (cont'd)

Formats: %d, %f %c, %x, others

```
int x=4;
```

```
char str [100];
```

```
FILE *f;
```

```
F = fopen ("myfile", "w");
```

```
fprintf (F, "%d", x);
```

```
sprintf (str, "%d %d %s", 12, x,  
         "hello"); => "12 4 hello"
```

FORMAT fields must match in type and #

Formatted Input

Formatted input and “input” strings

//Read from file

```
fscanf (FILE *in, const char *fmt,  
<ptr args ... allocated>);
```

// ”Read from” (i.e. parse) string

```
sscanf (const char*, const char *fmt,  
        <ptr args ... allocated >);
```

Also: scanf

Formatted Input (cont'd)

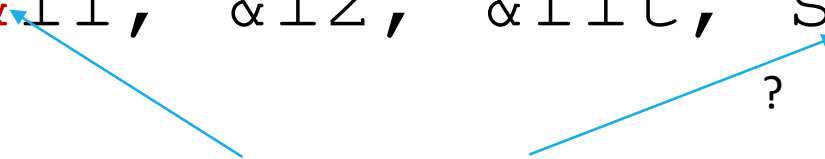
```
int i1, i2;
```

```
float flt;
```

```
char str1[10], str2[10];
```

```
char *str1, *str2; // ok?
```

```
sscanf ("11 12 34.07 keith ben",  
        "%d %d %f %s %s",  
        &i1, &i2, &flt, str1, str2);  
        ?
```



Do not forget to allocate and use &!

Project #1

Binary I/O

- We like ascii char-based files: `cat, more, emacs, ...`
- But they take up space
 - integer: `1234567890`
- Use binary I/O to read/write fewer bytes which saves space and is faster
- Works for fixed-size data items and lots of 'em
 - but must be memory-contiguous (i.e. an array)

```
fread/fwrite (void *buffer,  
              size_t size, size_t nitems,  
              FILE *f);
```

Binary I/O Example

```
typedef struct S{
    int ss;      // 8 digits
    int phone;   // 10 digits
} info_t;
info_t mine = {12345678, 5384937474};
// fopen F for write
fprintf (F, "%d %d", mine.ss, mine.phone);
fwrite ((void *)&mine, sizeof (info_t), 1, F);
```

ascii file: 12345678 538493747 = $(8+10)*1$ bytes

binary file: ^a93e&^%8 = $(4+4)*1$

```
fread ((void *)&mine, sizeof (info_t), 1, F);
printf ("%d %d\n", mine.ss, mine.phone);
```

suppose info_t mine [100000];

Random I/O

- Thusfar, all I/O was assumed to be sequential
 - `fgets, fgets, ...` returns consecutive lines
 - current **file ptr** advances sequentially
- Sometimes we need random I/O
- Think of any examples of applications that access data in a non-sequential pattern?


Random I/O

- Random I/O

- Advance file offset pointer w/o reading/writing

```
#include <stdio.h>
```

```
int fseek (FILE *stream, long off,  
           int whence);
```

 new offset in bytes

whence **of one** SEEK_SET, SEEK_CUR, SEEK_END

 from the beginning

// other call: what is the current file offset?

```
long ftell (FILE *stream);
```