

CSci 4061

Introduction to Operating Systems

Input/Output: Low-level cont'd

File descriptor magic

Chapter 4

File Descriptor Inheritance

FDs and FD table are inherited by children at `fork()` time!

```
int fd, fd1, fd2, pid;
fd = open ("my_file", O_RDONLY, 0);
pid = fork ();
if (pid != 0) {
    read (fd, ...);
    fd1 = open ("foo", ...);
}
else {
    read (fd, ...); // the same fd
    fd2 = open ("bar", ...);
    ...
}
```

Important for pipelines (later):

```
shell> cat foo.bar | grep jon
```

Re-direction: another key CS concept!

- What is the default input/output location?

```
shell> cat foo.bar
```

```
shell> cat foo.bar > baz.out
```

What is this? **Output redirection**

How does this work?

Let's build up to this

1. duplicating file descriptors
2. fd preservation

1. Duplicating File Descriptors

Sometimes you need to change what an `fd` points to (why? redirection)

```
#include <unistd.h>
```

```
int dup2 (int fd1, int fd2);
```

Closes `fd2` (if open) -> frees up `fd2`

Makes `fd2` now point to what `fd1` points to

Output Re-direction


```
shell> cat foo.bar > baz.out
```

What is happening? ... writes to stdout are redirected to `baz.out` ... but how?

Output redirection (`my_cat.c`)

```
fd = open ("baz.out", O_CREAT|O_WRONLY, ...);  
dup2 (fd, 1); // 1 now refers to fd  
close (fd);   // might as well ...  
write (1, ...); // writes to "standard out" => baz.out
```

Output Re-direction (cont'd)



Output Re-direction (cont'd)

Output redirection (my_cat.c)

```
fd = open ("baz.out", O_CREAT|O_WRONLY, ...);  
dup2 (fd, 1); //close stdout and 1 refer to fd  
close (fd);   // not needed but good style  
write (1, ...); // writes to "standard out" => baz.out
```

“lost” stdout - how can we re-open it?

```
fd = open ("/dev/tty", O_WRONLY);
```

Input Re-Direction

- How would we handle input redirection?

```
shell> wc -c
```

```
Hi there jon
```

```
^D
```

```
shell> 13
```

Suppose I want to use a file instead

```
shell> wc -c < mydata.txt
```


Input Re-Direction (cont'd)

my_wc.c:

```
fd = open ("mydata.txt", O_RDONLY, ...);  
dup2 (fd, 0);  
close (fd); // not needed but good style  
  
// reads from "standard in" directed to mydata.txt  
read (0, ...);
```

Re-Direction

- Re-directing within a single process not so useful, let's return to the shell
 - If I want to read/write to a file, just do it!
 - How does the shell actually do it for us?
- ```
shell> cat foo.bar > baz.out
```
- // cat delivers output to stdout and we don't  
// want to **modify** the code of cat or ls or ....

Step 2: fd preservation

# fd Preservation

```
int fd, pid;
char fd_str[2];

pid = fork ();
if (pid != 0) {
 read (fd, ...);
}
else {
 fd = open ("source_file", O_RDONLY, 0);
 read (fd, ...);
 sprintf (fd_str, "%d", fd);
 // fd's preserved through exec!
 execl (". /foo (char*)" "foo", (char*) fd_str);
 ...
}
```

```
foo.c
int main (int argc, char *argv[]) {
 int fd;
 fd = atoi (argv[2]);
 read (fd, ...); // will read from "source_file"
 ...
}
```

# I/O Redirection by the Shell

`cat foo > bar`

`wc < baz`

- Preservation of fd's
- Go back to `dup2`

# Inside `cat`

...

```
write (1, ...);
```

....

By default `cat` writes to `stdout`

This is true of virtually all shell commands/programs  
(read from `stdin`, write to `stdout`)

How do we get `cat` to write somewhere else?

# The Shell: dup + preservation

```
shell> myprog source > output
```

```
inside "myprog"
write (1, ...);
```

```
int pid;
pid = fork ();
if (pid != 0) { // shell parent
 ...
 wait (NULL);
}
else { // shell child
 // redirect name of file provided to the shell '>' case
 int o_fd;
 char *s_file; // extract from argv -- source
 char *o_file; // extract from argv -- output
 o_fd = open (o_file, O_WRONLY);
 dup2 (o_fd, 1);
 exec1 ("myprog ...", (char*)&s_file, NULL);
}
```