

CSci 4061

Introduction to Operating Systems

File Systems: Basics

Chapter 5

File as Abstraction

- Container for related information
- Named
- Associated attributes
- Persistent


Naming a File

```
creat/open ("path/name", ...);
```

Links: files with multiple names

Each name is an alias

```
#include <unistd.h>
int link (const char *original_path,
          const char *new_path)
```



cannot exist as
a file already

```
link ("foo", "bar"); // "bar" refers to file "foo"
unlink ("bar");      // remove name "bar"
// if file is open by someone, will not actually get deleted until
// all fd's to it are closed
```

File Attributes: Access to metadata

```
#include <sys/stat.h>
int fstat (int filedes, struct stat *buf)

int stat (const char *pathname,
          struct stat *buf)
```

Structure contains file/directory info:

```
off_t st_size;      // file size
nlink_t st_nlink;   // links
mode_t st_mode;     // type + permission
time_t st_mtime;    // last modification time
```

`fcntl` can also be used to set or get lower-level attrs

Exercise: Metadata

- Write a program that monitors a given file every minute `[sleep(60)]` and if the file size has changed, it outputs the new size to stdout

```
#include <sys/stat.h>
int stat (const char *pathname,
          struct stat *buf)
```

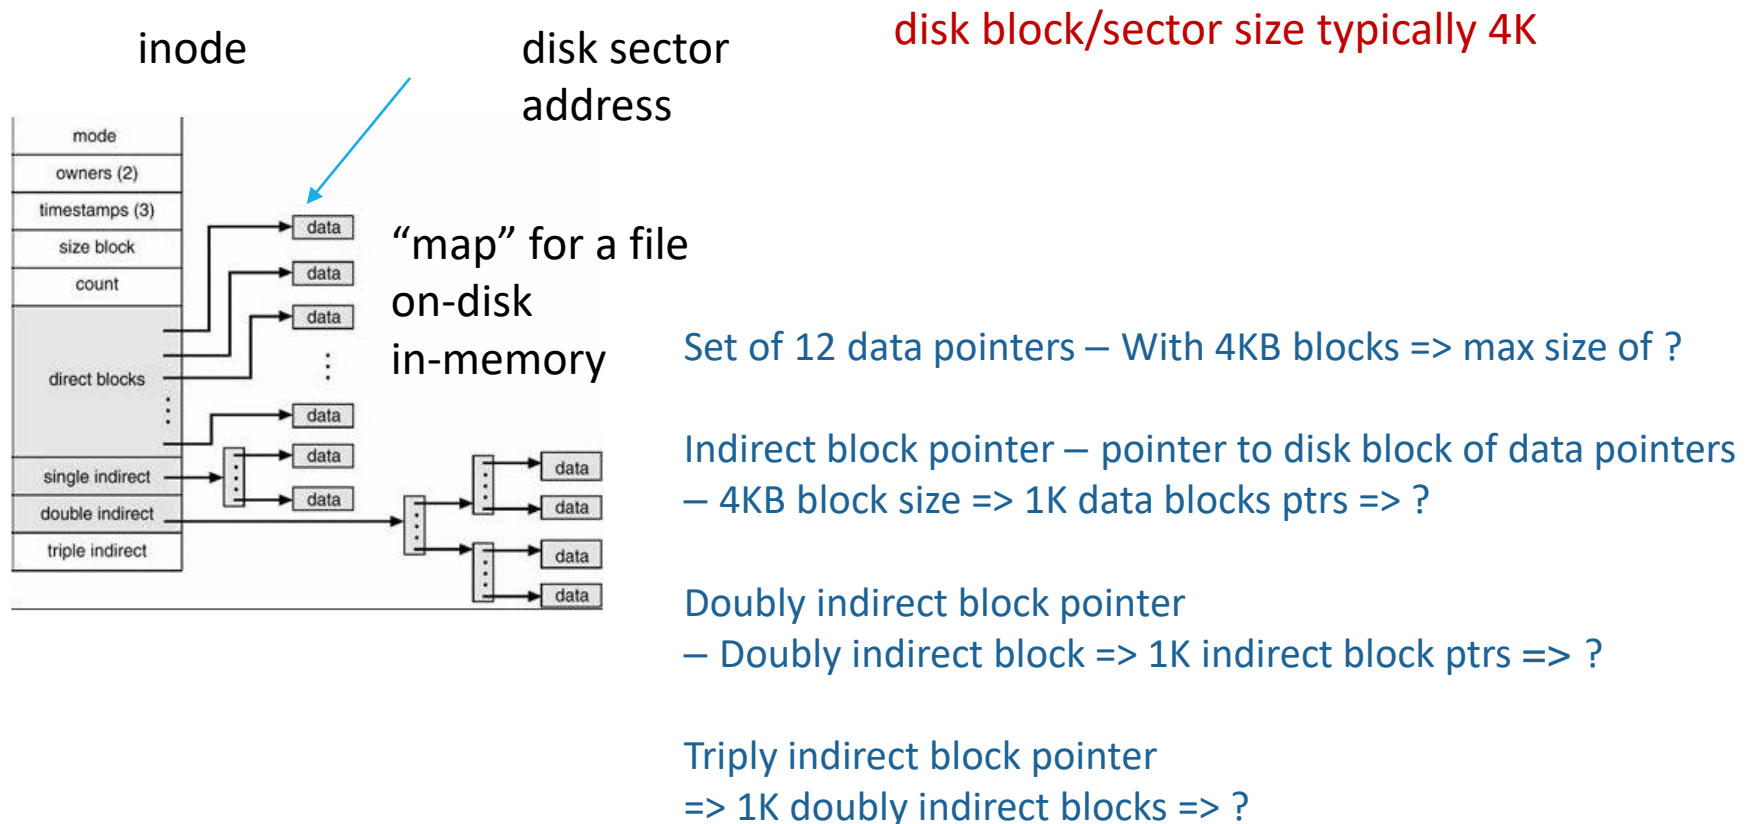
Structure contains file/directory level info:

```
off_t st_size;           // file size
nlink_t st_nlink;        // links
...
```

Example

```
void mon_size (char *fname) {  
    struct stat sb;  
    off_t psize = -99;  
  
    while (1) {  
  
        sleep (60);  
    }  
}
```

Storing File Meta-data: Unix inode



Filesystem

- Directory is a file as well
 - it has an inode
 - what are file contents?
- Filesystem
 - Files
 - Directories
 - Free disk sectors (free list)
 - Root dir

Filesystem (cont'd)

- On-disk organization
 - inode for root dir of filesystem “/” stored in well-known sector on the disk
 - inode for disk sector *free-list* also stored in a well-known sector on the disk
 - inode table or file (inode #, sector)
 - These are stored in the *superblock*

Unix file types/modes

- Indicated by the first character in `ls -l`
 - - regular file
 - d directory
 - c character special file
 - b block special file
 - p pipe
 - s socket
 - l symbolic link

File types

- Within `stat` structure:

```
struct_t stat st;  
stat ("foo", &st);
```

Macros:

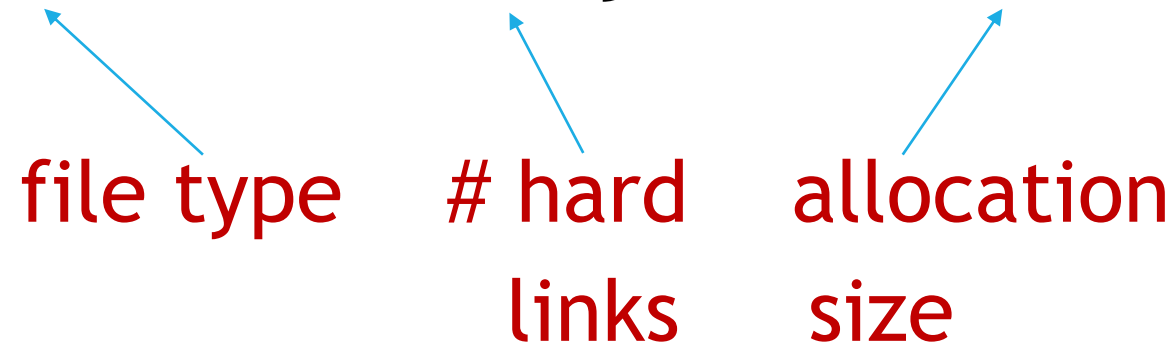
```
int S_ISDIR (st.st_mode);  
int S_ISREG (st.st_mode);  
int S_ISSOCK (st.st_mode);  
...
```

see P. 158 (Table 5.1)

Another look at ls -l

Example:

```
drwx-xr-x  3  jon  fac 4066 Nov 2 09:14 st
```



file type # hard links allocation size

File permissions

Operations (**r**, **w**, **x**): **r** read, **w** write, **e**xecute

Subjects (**u**: **u**ser/owner, **g**: **g**roup, **o**: **o**thers)

Users may belong to any number of groups
(type `groups` at the shell)

File permissions (cont'd)

```
shell> ls -l
```

```
drwxr-xr-w    3  jon  fac  46
```

u g o owner group

When a file is created it is given a restricted permission and a default group

You can broaden or further restrict permissions

File permissions (cont'd)

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
int chmod (char *path, mode_t mode);
```

0077

each octal digit is rwx (or 1) for ugo

000, 111, 111



prefer symbolic flags: man fstat, e.g. S_IRGRP : GRouP has Read)

Also at the command-line (absolute)

```
chmod 0077 st.txt
```

Also at the command-line (relative)

```
chmod go-xr st.txt
```

```
chmod u+xrw st.txt
```


Filesystem semantics: Unix

- Two processes open the same file
- Reader sees most recent write
- One reader and one writer - run together
 - File “foo” contains
“aaaaaaaaaaaaaaaaaaaaaaaaa”

Filesystem semantics (cont'd)

```
// reader.c
```

```
#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
```

```
void main () {
    int fd, n;

    char c, buf[100];
    read (0, &c, 1);
    fd = open ("foo", O_RDONLY);
    n = read (fd, buf, 10);
    buf[n] = '\0';
    printf ("buf=%s\n", buf);
    ...
}
```

```
// writer.c
```

```
#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
```

```
void main () {
    int fd, n;

    char c, buf[100] = "bbbbbbbbbbbbb...";
    read (0, &c, 1);
    fd = open ("foo", O_WRONLY);
    write (fd, buf, 10);
    read (0, &c, 1);
    close (fd) }
```

Power of IDs

- Real user-id: user that actually initiated a process
 - Not executable owner!
- ```
-r-x ... 1 jon fac 203 Feb 10 10:47 test
bill> /usr/jon/test
```
- Effective user-id: user that system associates with the process for purposes of protection
    - Usually the same as the real user-id: this would be?
    - Sometimes want effective user-id to that of the file owner and not the user ... why?

# Power of IDs (cont'd)

- How do to it?

```
jon> chmod u+s test
```

```
-r-s... 1 jon fac 203 Feb 10 10:47 test
```

```
bill> /usr/jon/test
```

has privilege of 'jon'

```
jon> chmod g+s
```

has privilege of group 'fac'

# Masks

```
creat ("my_file", 0777);
```

- Expectation:

```
shell> ls -l
-rwxrwxrwx jon ... my_file
```

- Instead:

```
-rw----- my_file
```

- What happened?

- To prevent against accidental exposure, Unix sets a default mask with your process (type `umask`)
- Typically: 077 (1 means mask out)

```
creat ("name", PERM & (~mask)); AND
regular files also mask out execute
```

# Masks (cont'd)

```
creat ("name", PERM & (~mask));
```

umask is 022: what is this one?

To change the mask:

```
mode_t umask (mode_t newmask);
```