

CSci 4061



Introduction to Operating Systems

Module 4: Communication
IPC: Basics, Pipes
Chap 6-6.2

Communication

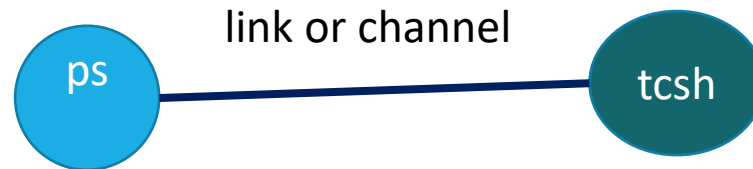
- Abstraction: conduit for data exchange between two or more processes (or threads)

IPC in Unix

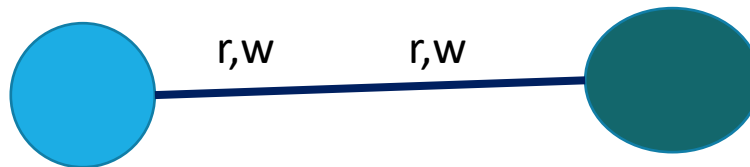
- Pipes: most basic form of IPC in Unix
 - process-process
 - `ps -u weiss039 | grep tcsh // what happens?`
- Pipe has a “read-end” (receive) and a “write-end” (send) : think of this actually as a 
 - FIFO communication
 - (write A, write B, read->A, read->B)
 - “Bi-directional”



IPC in Unix (cont'd)




- Pipe allows communication between a **parent and child or related processes**



Pipes

```
#include <unistd.h>  
int pipe (int ends[2]); // returns -1 on failure
```



output parameter

`ends` is a 2-integer `fd` array that represents the ends of the pipe

`ends[0]` is the “read-end” (receive) and
`ends[1]` is the “write-end” (send)

Integrated into filesystem

Link is “named” by the pipe but we do not name the reader/writer processes

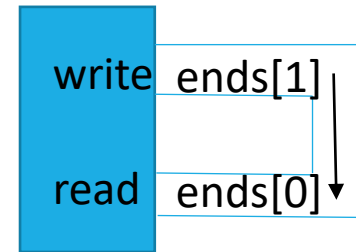
How can `pipe` fail?

Pipes and FD

Simple pipe example: single process

```
#include <unistd.h>
#include <stdio.h>
#define MSGSIZE 16
char *msg1 = "hello, world #1";
char *msg2 = "hello, world #2";
void main () {
    char inbuf [MSGSIZE];
    int ends[2];

    if (pipe(ends) == -1) {
        perror ("pipe error");
        exit (1);
    }
    ...
```



Simple pipe example (cont'd)

// write (send) down pipe

```
write (ends[1], msg1, MSGSIZE);
```

```
write (ends[1], msg2, MSGSIZE);
```

// read (receive) from pipe

```
read (ends[0], inbuf, MSGSIZE);
```

```
fprintf (stderr, "%s\n", inbuf);
```

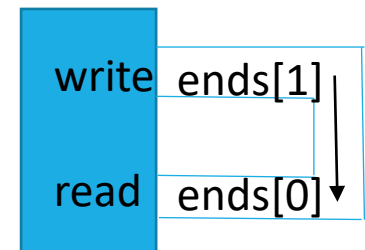
```
read (ends[0], inbuf, MSGSIZE);
```

```
fprintf (stderr, "%s\n", inbuf);
```

Output is:

hello, world #1

hello, world #2



Read and write

```
write (ends[1], msg, MSGSIZE);  
read  (ends[0], inbuf, MSGSIZE);
```

Read may not get everything but it “usually does” up to *max* (MSGSIZE, and pipe contents)

blocks if pipe is empty

Pipe have finite size (e.g. 4K/8K)

write blocks if not enough space

why is there a limit?

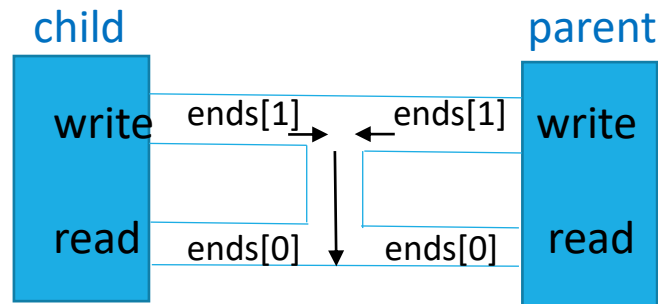
Simple pipe example: multi process

```
void main () {  
    char inbuf [MSGSIZE];  
    int ends[2], j;  
    pid_t pid;  
  
    if (pipe(ends) == -1) {  
        perror ("pipe error");  
        exit (1);  
    }
```

Multi process (cont'd)

```
pid = fork ();  
if (pid == 0) {                                // child sends into pipe  
    write (ends[1], msg1, MSGSIZE);  
    write (ends[1], msg2, MSGSIZE);  
}  
else if (pid > 0) {                            // parent receives from pipe  
    read (ends[0], inbuf, MSGSIZE);  
    fprintf (stderr, "%s\n", inbuf);  
    read (ends[0], inbuf, MSGSIZE);  
    fprintf (stderr, "%s\n", inbuf);  
    wait (NULL);  
}
```

why does this work across
these processes?



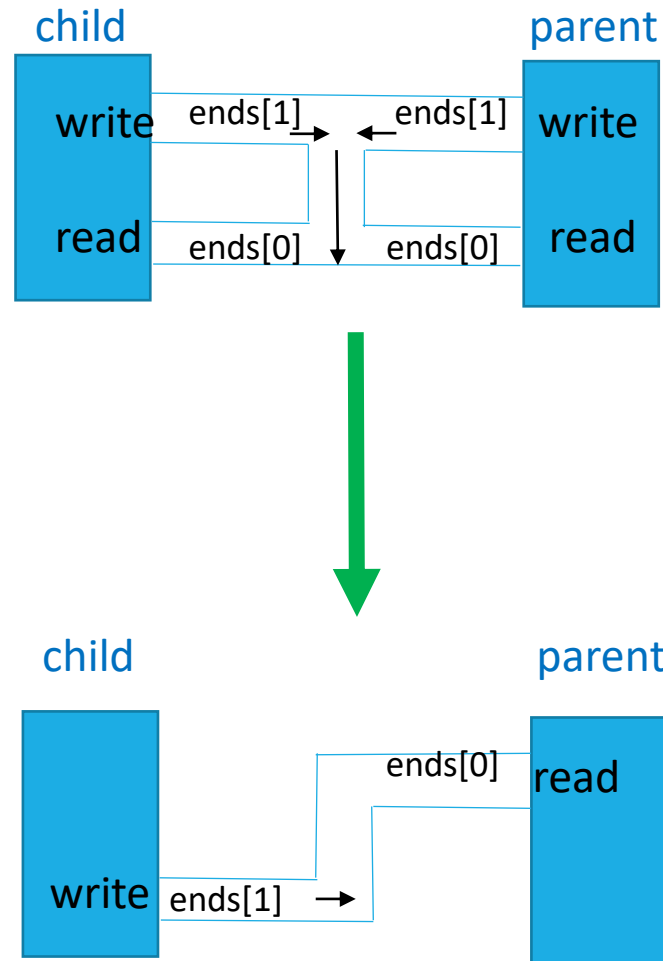
Issues

- Potential problem
 - If both processes write into the pipe, what would happen?
 - Usually, one writes and other reads
 - Either way, writes would likely be serialized

Resolving

```
if (pid == 0) {                                // child sends into pipe
    close (ends[0]);
    write (ends[1], msg1, MSGSIZE);
    write (ends[1], msg2, MSGSIZE);
}
else if (pid > 0) {                            // parent receives from pipe
    close (ends[1]);
    read (ends[0], inbuf, MSGSIZE);
    fprintf (stderr, "%s\n", inbuf);
    read (ends[0], inbuf, MSGSIZE);
    fprintf (stderr, "%s\n", inbuf);
    wait (NULL);
}
```

New picture



Typical Pipe Use Case

- Near infinite stream of data from producer to consumer
 - consumer (reader) had better keep up with producer (writer) and vice-versa
 - why?
 - `cat * | wc`
- You cannot `fseek/seek` a pipe `fd`

More on pipes

- `close` write-end (**and** no processes have pipe open for write) **and** pipe is empty:
 - `read` returns a 0
- `close` read-end and write-end is open:
 - `write` kills the process!
 - “broken pipe”
- Pipes are limited to parent-child siblings, related process relationships must share `fds`

Example: Knock-Knock

- <https://www.youtube.com/watch?v=uky0JpQzDNI>

Example: Knock-Knock

- Protocol: sequence of messages

P: “k-k” “orange” “aren’t you glad this isn’t Java?”
C: “w-t?” “orange-who?”

```
graph LR; P1["P: 'k-k'"] --> C1["C: 'w-t?'"]; C1 --> P2["P: 'orange'"]; P2 --> C2["C: 'orange-who?'"]; C2 --> P3["P: 'aren't you glad this isn't Java?'"];
```

P: w, r, w, r, ..

C: r, w, r, w, ...

Solution

Sol 1:

```
pipe (ends);
```

```
fork ();
```

```
// parent
```

```
write (ends[1], "k-k", ...)
```

```
read (ends[0], buf, ...);
```

```
write (ends[1], "orange", ...);
```

```
read (ends [0], buf, ...);
```

```
write (ends[1], "aren't ...", ...);
```

```
// child
```

```
read (ends[0], buf, ...);
```

```
write (ends[1], "w-t?", ...);
```

```
read (ends [0], buf, ...);
```

```
write (ends[1], "orange-who?", ...);
```

```
read (ends [0], buf, ...);
```

Issues?

Better one?

Solution

Sol 2:

```
int P_C[2], C_P[2];  
pipe (P_C);  
pipe (C_P);  
fork ();
```

// parent

```
close (P_C[0]);  
close (C_P[1]);  
write (P_C[1], ...)  
read (C_P[0], ...);  
...
```

// child

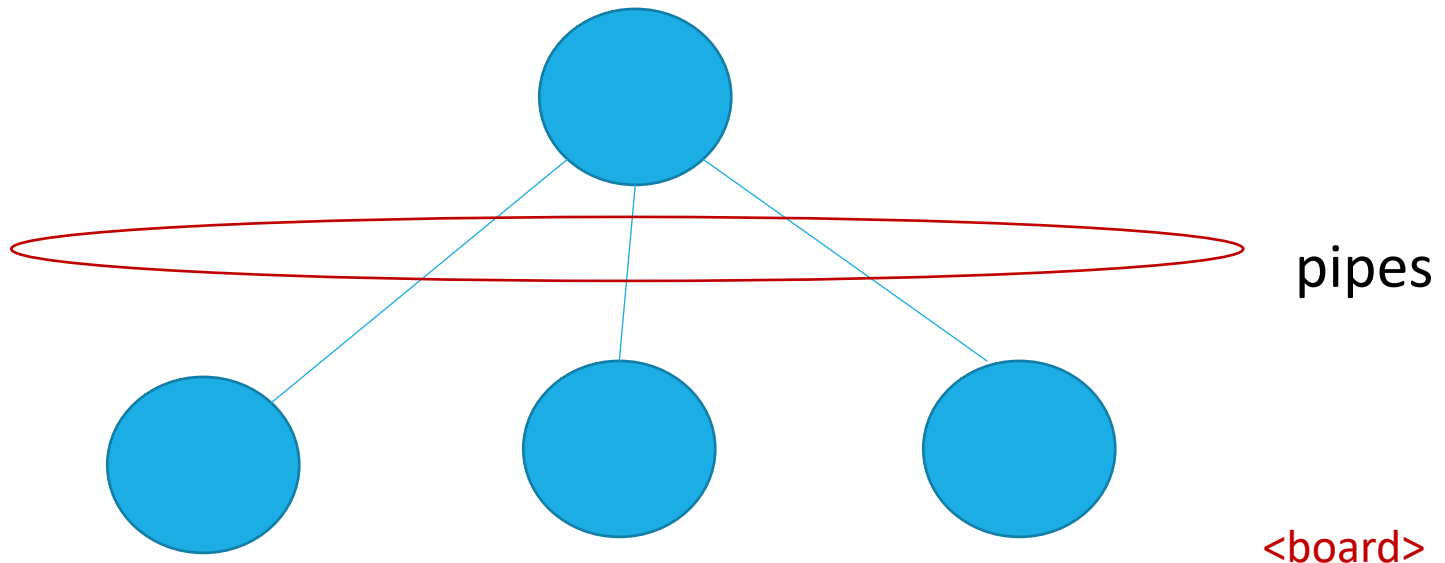
```
close (P_C[1]);  
close (C_P[0]);  
read (P_C[0], ...);  
write (C_P[1], ...);  
...
```

Takeaway Lesson

- Knock-knock example ...
- Need two-way communication
- Most likely will need a pair of pipes

Non-blocking pipes: example

- Children may inform the parents of various events or ask for things to do ... BUT this is unpredictable ...
- Suppose we do blocking I/O?



Non-blocking pipes

- Default I/O behavior is blocking
- Non-blocking I/O can be handy
- Since pipe is a file ... can control attributes

```
#include <fcntl.h>

int fcntl (int fd, int cmd, ...);
int ends[2], flags, nread;

pipe (ends);
flags = fcntl (fd, F_GETFL, 0);
fcntl (ends[0], F_SETFL, flags | O_NONBLOCK);
...
nread = read (ends[0], buf, size);

// if nothing to read, returns -1, errno set to EAGAIN
```


Pipes in the shell

- `ps -u weiss039 | grep tcsh`
- How does the shell do it?

```
pipe (ends);  
if (childpid = fork ()) == 0) {  
    dup2 (ends[1], 1);  
    close (ends[0]);  
    execl ("/bin/ps", ...);  
}  
else {  
    dup2 (ends[0], 0);  
    close (ends[1]);  
    execl ("/bin/grep", ...);  
}
```

Sending Discrete “Data”

- Sending a message into a pipe or any fd

```
typedef struct {  
    int x;  
    int y;  
    char str[20];  
} message_t;
```

Must be contiguous data

```
message_t m1, m2;  
int ends[2];
```

```
pipe (ends); // check for error!
```

```
// send m1 into the pipe
```

```
write (ends[1], &m1, sizeof (message_t));
```

```
// pull data into m2 from the pipe
```

```
read (ends[0], &m2, sizeof (message_t));
```