

# CSci 4061


# Introduction to Operating Systems

Module 5: Threads

**(Thread-Basics)**


Chapter 12

# Threads




- Abstraction: for an executing instruction stream
- Threads exist within a process and **share its resources** (i.e. memory)
- But, thread has its own stack and “PC”
- Default: always 1 thread (implicit)

# Another View



# Two Threads Sharing a CPU




What may cause a switch?

# Thread Benefits

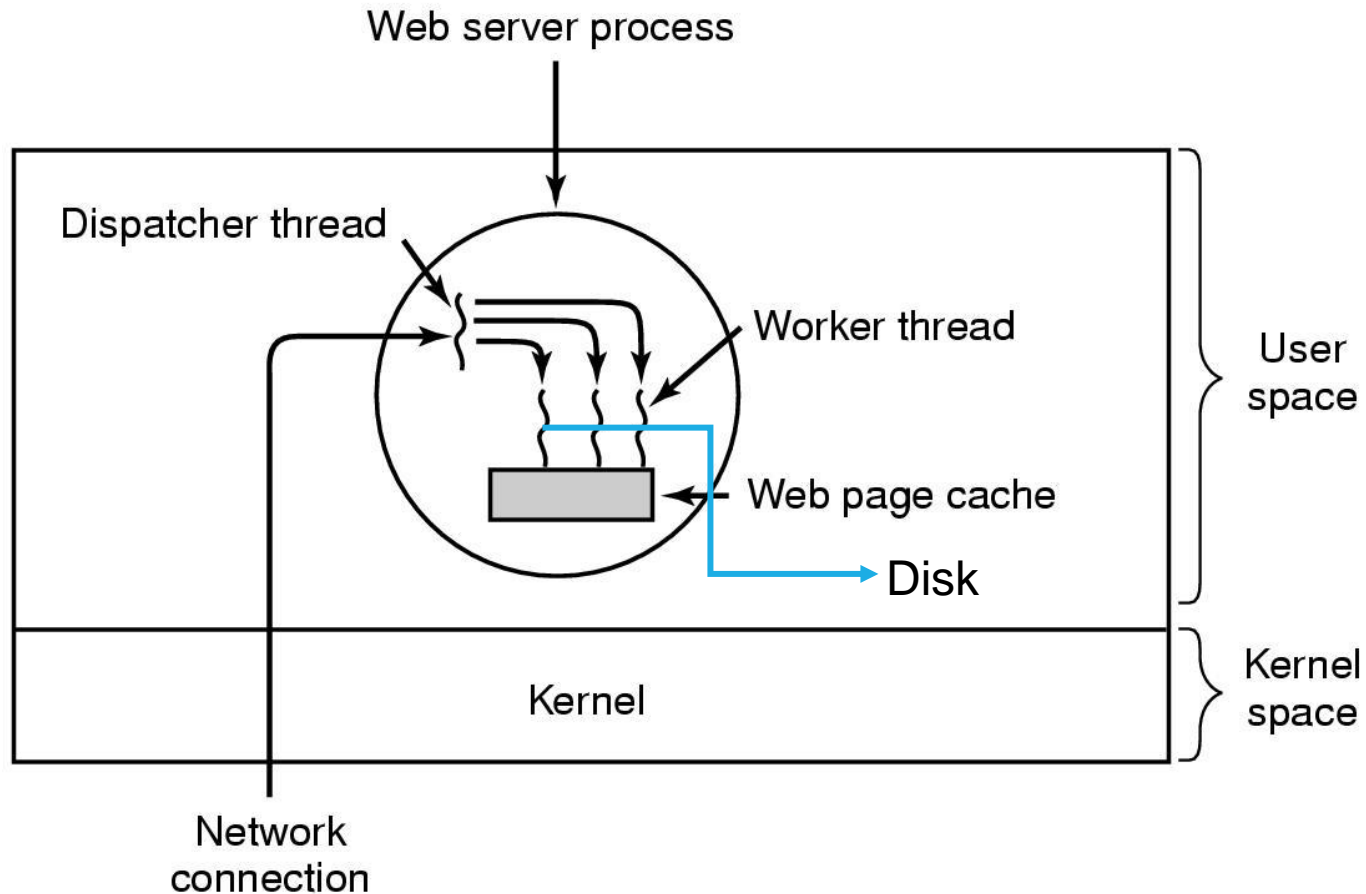
- Concurrency
  - when one blocks, another runs
- Modularity
  - decompose functionality
- Parallelism
  - threads running in parallel - multi-core
- Scale
  - more threads than processors/processes
- Overhead
  - cheaper than processes

# Threads Example: editor



When one blocks, another can run ...

# Thread example: Web server...



When one blocks, another can run ...

# How to Quantify the Benefit?

- Web request: read request from network, fetch page, write request to network
- $T_{\text{req}} = T_{\text{read}} + T_{\text{serve}} + T_{\text{write}}$
- $T_{\text{req}} = T_{\text{read}} + h * T_{\text{cache}} + (1-h) * T_{\text{disk}} + T_{\text{write}}$
- Can we improve the performance of a single request with threads?
- Can we improve the performance of the “service”, requests/time?



# Thread Example: Web server

dispatcher (...) {

```
while (TRUE) {  
    // 1. read (req ~ URL)  
    get_next_request (&req);  
    handoff_work (&req, &buf);  
}
```

worker (...) {

```
wait_for_work (&buf, &req)  
// 2. service  
look_for_page_in_cache (&req, &answer);  
if (page_not_in_cache (&answer) {  
    read_page_from_disk (&req, &answer);  
    put_page_in_cache (&req, &answer);  
}  
// 3. write  
return (&answer);
```

- How are these threads interacting?
  - Shared memory: threads share buffer, cache
  - Threads share globals, heap, NOT stack

Looks great

- Drawbacks?
- Alternatives?

# Drawbacks

- Sharing
  - Synchronization is needed to protect shared data structures: Web server? Editor?
  - Assume: threads may be switched unpredictably!
  - Failure: no isolation
- Thread-safety (related to Sharing)
  - Not all system calls may be thread-safe
  - System call (or any call) that can be executed concurrently by multiple threads
- Global variables
  - Per thread globals may be needed

# Drawbacks: Sharing/Thread-safe

```
int counter = 0;
int increment_counter () {
    counter ++; // counter = counter +1
return counter;
}
```

problem? Suppose threads **T1** and **T2** call it

# Unrolling counter = counter + 1

load counter->R

incr R

store R->counter

# Thread Safety (cont'd)

```
int counter = 0
lock_type counter_lock;
int increment_counter () {
    // lock is held or free: if held, caller is blocked
    lock (counter_lock);
    counter++;
    unlock (counter_lock);
    return counter;
}
```

# Locks

- Just to be safe, shouldn't I always put locks around my code?
  - Locks reduce concurrency and performance, use only when needed!

# Drawbacks:per thread globals

T1

...

syscall\_1

sets errno

... {T1 switches to T2}

reads errno

T2

...

syscall\_2

sets errno

...



- In Unix, *errno* is a global variable in shared library
- Options to guarantee error reporting is thread-safe?
  - Use locks
  - Eliminate global: return error code
  - Define *errno* “service” or macro

```
#define errno _special_thread_errno (thread_id)
```



# Alternatives to Threads

- Want concurrency
  - Goal: If a program (or part of a program) cannot make progress due to blocking, then allow:  
some other part of the program to make progress
- Options?