



CSci 4061

Introduction to Operating Systems

Module 5: Threads (Implementation/Models)


Chapter 12

Two Threads Sharing a CPU



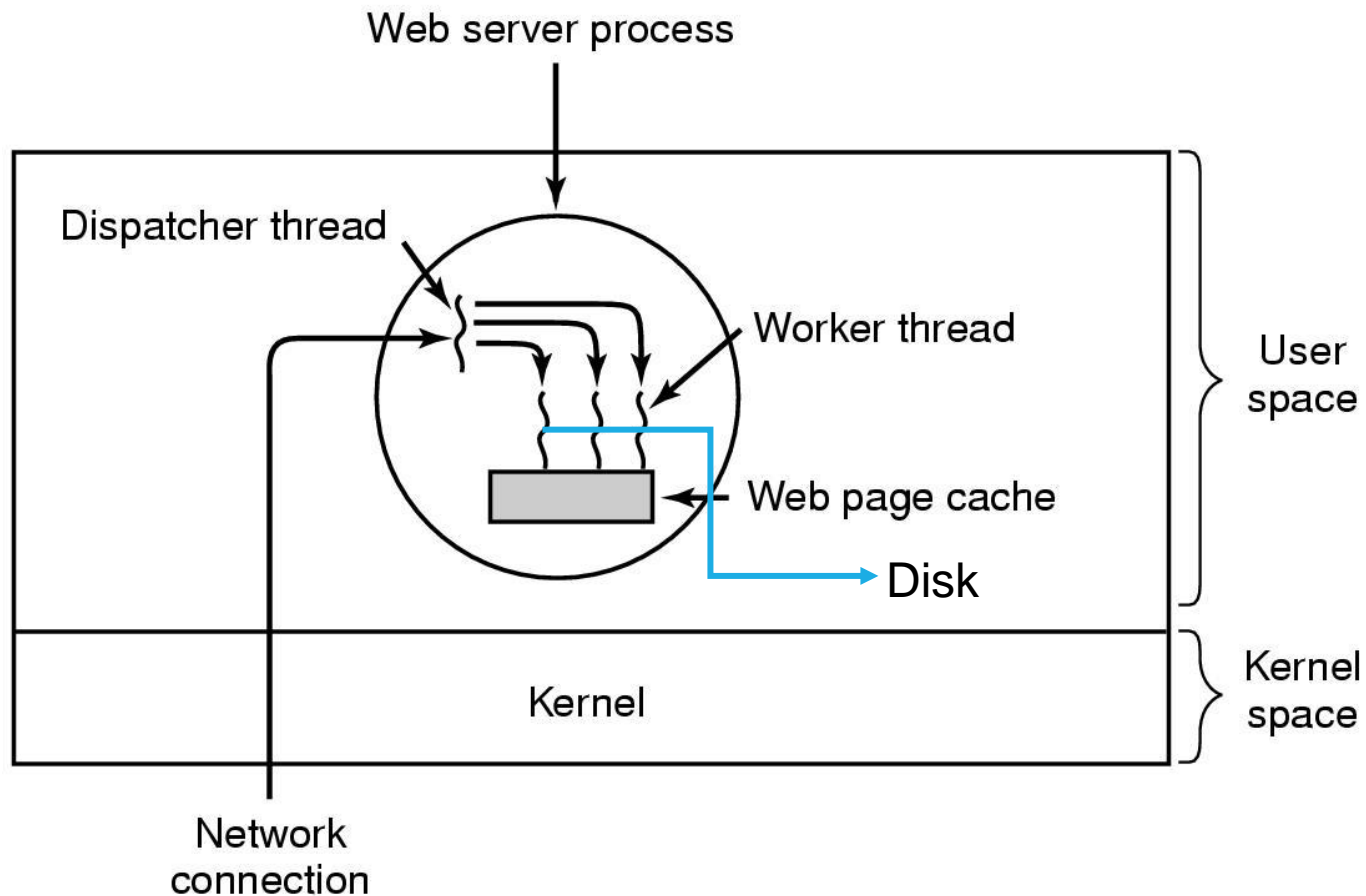
What may cause a switch?

Threads Example: editor



When one blocks, another can run ...

Thread example: Web server...




When one blocks, another can run ...


Thread Overview

- Thread advantages
 - modularity, concurrency
 - sharing, cheap
- Sharing is a double edged sword
 - race conditions, failure
- **Implementation models**
 - user, kernel, hybrid
- **Programming models**
 - dispatcher and team

Implementing Threads in User Space



Implementing Threads in the Kernel



A threads package managed by the kernel

User vs. Kernel Threads

- User thread advantages
 - no thread system calls! --cheaper
 - more scalable
 - more portable
 - custom control and scheduling
 - blocking is a big problem!

```
dispatcher(...) {  
    while (TRUE) {  
        get_next_request (&req);  
        handoff_work (&req, &buf);  
    }  
}
```


```
worker(...) {  
    wait_for_work (&buf, &req);  
    look_for_page_in_cache (&req, &answer);  
    if (page_not_in_cache (&answer) {  
        read_page_from_disk (&req, &answer);  
        put_page_in_cache (&req, &answer);  
    }  
    return (&answer);  
}
```


Kernel Threads

- Advantages
 - thread can block: OS can pick another from same process
 - can exploit multiprocessors
- Hyper-threading
 - hardware support for threads!

Hybrid Thread Models

- One-To-One (Linux, WinXY, Posix)
 - each user-level thread maps to a kernel thread
 - if on a multicore system




User thread

Kernel thread

Hybrid Thread Models

- Many-to-One (Java)
 - Many user-level threads mapped to a single kernel thread



How do I program them?


General Thread Operations

- `Create/Fork`
 - Allocate memory for stack, perform bookkeeping
 - Parent thread creates child threads
 - Associates function to execute
 - Returns an id
- `Destroy/Cancel`
 - Release memory (or recycle), perform bookkeeping
- `Suspend (e.g. Sleep) -> blocked`
 - Resume->unblock, Yield->deschedule
- `Wait/Join`
 - Wait for something, e.g. child finishing

Inside Threads

- A thread contains
 - pc
 - sp
 - registers
 - child threads
 - state
- What about open files?

Threads in Action



- `main` starts executing creates T1 and T2
- `main` blocks at `wait`
- Switches to T1 or T2, say T1
 - blocks at `read`
- Switches to T2

Thread Models

- **Dispatcher-worker (master-slave)**

- a master process/thread receives request for work
- generates/dispatches a thread to service work request
 - e.g. threaded server

- **Two options:**

- master can create a thread on “as needed basis” *pop-up*
- master can keep a thread pool
 - may reduce perceived latency of creating threads to service request
 - issues?

Thread pop up vs. pool

- Pop up
 - ^ latency of serving request
 - can size the # of threads to workload
 - what to do when workload is very high?
- Thread pool
 - how to size the pool?
 - too many threads may waste resources
increase context switching
 - too few threads increase latency
 - could make the pool dynamic!

Thread Models (cont'd)

- **Team**

- a collection of peer threads working on some part of a problem together
- identical threads (parallelism):
 - parallel program running on shared-memory multiprocessor, multi-core
 - n threads are created and each are given a share of the problem e.g. scale element of a matrix
- different threads (concurrency):
 - editor example

Thought Question

- On a uniprocessor
 - Threaded matrix multiply program
 - NxN matrix (N is large) and sitting in memory
 - Create 4 threads each responsible for $\frac{1}{4}$ of the matrix multiply operations
 - Time the 4-threaded version and compare with a single threaded version
 - The 4-threaded version does worse---WHY?

Remember

- System may time-slice your thread
- You should assume that a thread could be switched at any time ... your program should still work
 - A program that fails 1 out of 10^{100} runs is buggy
- This will make your code much more portable

Thread Safety

- Two or more threads call code that has the potential for race conditions.
- Deal with it.
- Linux man pages will tell you if a syscall is thread-safe...or not