

# CSci 4061

## Introduction to Operating Systems

### Synchronization Basics: Locks

# Motivation


- Issues
  - Threads communicate through shared variables
  - K ready or “runnable” threads => can’t predict which one is running at any particular time

# Synchronization Outline


- Basics
- Locks
- Condition Variables
- Semaphores - if time
- Issues

# Basics

- Race condition: threads + shared data
- Outcome (data values) depends on who gets there first/last



- Possible values for `i` at the end of execution? 7,8,4,5!
- Shared variables = heap, globals, within the process
- Races => inconsistency or errors




- If buffer is nearly full=> may overwrite or overflow

# Problem

- Problem: we have limited control on when threads will run
- Need: orderly execution or cooperation
- Solution: synchronization
- Real life: washing dishes
  - Wash then dry
  - No two people washing at the same time

# Synchronization

- Constrain the set of interleavings
  - Can't prevent scheduler from switching them out
  - But threads can stay out of each others way



- Critical section
  - Region of code where shared access may lead to races
  - Constrain access to critical section
  - Only 1 thread at a time in the critical section

# Critical section: How to do it?

- Threads **voluntarily** spin or block (wait) if another is in the critical section

entry      =>      possibly block or spin  
    <CS>                      <CS>  
exit

- Examples of critical section


```
if (free_buffer)
    insert_item
```

```
if i == 0
    => i = 5
else i = 7
```

```
if i == 0
    => i = 4
else i = 8
```

# How to identify a CS: good question!

- Black art
- Conservative (too big) ?
- Too small =?
- Mutual exclusion: simplest type of synch
  - Only 1 thread allowed in CS
  - CS is “atomic” (all or nothing)—can be interrupted, but no one else can get in
- Exit
  - Crucial to make it work!





# Related Issues

- Synchronization
  - Prevent bad things from happening
  - “wash then dry”, “no two washers...” (washing is a CS)
- Deadlock
  - Extreme case (misuse) of synchronization, everyone is blocked: join (self)
- Livelock
  - Everyone can run (not blocked) but no one can make progress
  - “one step forward, one step back”

# Synchronization construct for mutual exclusion (ME)


- Locks:
  - Object in shared memory
  - Operations: `acquire (lock)`, `release (unlock)`
  - Try to acquire a “held” lock => prevented
  - `acquire` lock before entering CS
  - `release` lock before leaving CS

```
Lock L;  
acquire (L);  
<CS>  
release (L);
```

Lock is EXPLICIT—have to use it correctly!

```
T1  
acquire (L);  
access to var X  
release (L);
```

```
T2  
access X // this is allowed!
```



# Use it carefully

T1

acquire (L1);  
access to var X  
release (L1);

T2

acquire (L2);  
access to var X  
release (L2);

# Inside a Lock

- Lock

```
boolean held;  
queue waiting_threads;
```

Mutex:

If held: acquire blocks thread and puts in on queue

If queue is non-empty: release removes a thread from queue and makes it unblocked

# Synchronization in Posix

- Posix mutex

```
#include <pthread.h>
pthread_mutex_t mtx = PTHREAD_MUTEX_INITIALIZER; // unlocked
```

//acquire

```
int pthread_mutex_lock (pthread_mutex_t *mutex);
```

//release

```
int pthread_mutex_unlock (pthread_mutex_t *mutex);
```

//return 0 on success, non-0 error code otherwise

```
gcc -o myProg myProg.c -lpthread
```

# Two Locks Deadlock; 2 Threads

- Deadlock - every thread is blocked

Lock L1, L2;

T1:

Acquire L1;

Acquire L2;

Release L2;

Release L1;

T2:

Acquire L2;

Acquire L1;

Release L1;

Release L2;

# Mutex Example

```
account act; // global shared state
```

```
// some number of deposit threads will be created
```

```
pthread_create (&t1, NULL, depositer, ...);
```

```
pthread_create (&t2, NULL, depositer, ...);
```

```
void *depositer (void *arg) {
```

```
    amount_t amt, val;
```

```
    //determine amt somehow
```

```
    ...
```

```
    val = deposit (&act, amt);
```

```
    ...
```

```
}
```

# Mutex example (cont'd)

```
pthread_mutex_t acc_mtx =  
    PTHREAD_MUTEX_INITIALIZER;
```

```
amount_t deposit (account *act,  
                  amount_t amount)
```

```
{  
    amount_t result;  
    pthread_mutex_lock (&acc_mtx);  
    act->balance += amount;  
    result = act->balance;  
    pthread_mutex_unlock (&acc_mtx);  
    return result;  
}
```

two threads  
calling deposit



# Thread safety

Suppose you are not sure a library call is thread-safe?

`rand ()` –

what can you do?

# Randsafe Example

```
#include <pthread.h>
#include <stdlib.h>

int randsafe (double *ramp) {
    static pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
    int error;

    pthread_mutex_lock (&lock);
    *ramp = (rand() + 0.5) / (RAND_MAX + 1.0);
    pthread_mutex_unlock (&lock);
    return;
```

# Are locks themselves safe?

- Yes!
- Must be possible for threads to concurrently call lock and unlock!
- All lock code is thread-safe

# Posix mutex (cont'd)

- Can test if lock is held

```
#include <pthread.h>
```

```
int pthread_mutex_trylock
```

```
(pthread_mutex_t *mtx)
```

- returns EBUSY if mtx is held

- Be careful: why?

```
if (pthread_mutex_trylock (&mtx) != EBUSY)  
    pthread_mutex_lock (&mtx);
```

# Posix mutex: Bounded Buffer

Need ME, why?

```
item_t remove_item (buffer *b) {  
    item t st;  
    if (b->next_slot_to_retrieve ==  
        b->next_slot_to_store) return ERROR;  
    st = b->items [b->next_slot_to_retrieve];  
    b->next_slot_to_retrieve++;  
    // adjust next_slot_store if needed  
    return st;  
}
```

# Posix mutex: Bounded Buffer

## Need ME:

```
pthread_mutex_t mtx = PTHREAD_MUTEX_INITIALIZER;
item_t remove_item (buffer *b) {
    item_t st;
    pthread_mutex_lock (&mtx);
    if (b->next_slot_to_retrieve ==
        b->next_slot_to_store) return ERROR;
    st = b->items [b->next_slot_to_retrieve];
    b->next_slot_to_retrieve++;
    // adjust next_slot_store if needed

    pthread_mutex_lock (&mtx);
    return st;
}
```

# Synchronization

- Mutual exclusion (ME) solved with locks
  - just have to use them correctly
- Want other kinds of synchronization

# Posix mutex (cont'd)

- Locks are limited to protecting shared variables only ... and they are **unconditional**
- Want richer synchronization

```
pthread_mutex_t mtx = PTHREAD_MUTEX_INITIALIZER;
item_t remove_item (buffer *b) {
    item_t st;
    pthread_mutex_lock (&mtx);
    if (b->next_slot_to_retrieve ==
        b->next_slot_to_store) return ERROR;

    st = b->items [b->next_slot_to_retrieve];
    b->next_slot_to_retrieve++;
    // adjust next_slot_store if needed
    pthread_mutex_lock (&mtx);
    return st;
}
```



# Posix mutex (cont'd)

- Locks are limited to protecting shared variables only ... and they are **unconditional**
- **Want richer synchronization**

```
pthread_mutex_t mtx = PTHREAD_MUTEX_INITIALIZER;
item_t remove_item (buffer *b) {
    item_t st;
    pthread_mutex_lock (&mtx);
    if (b->next_slot_to_retrieve ==
        b->next_slot_to_store) return ERROR; // block

    st = b->items [b->next_slot_to_retrieve];
    b->next_slot_to_retrieve++;
    // adjust next_slot_store if needed
    pthread_mutex_unlock (&mtx);
    return st;
}
```

# Need Richer Synchronization: ~ **conditional** synchronization

- Want producer (and consumer) to **conditionally** block if buffer full/empty

// should block if empty

```
item = remove_item (&b);
```

// should block if full

```
insert_item (&b, item);
```