

CSci 4061

Introduction to Operating Systems

CVs in Posix

Posix condition variables

```
#include <pthread.h>
```

```
pthread_cond_t cond =  
    PTHREAD_COND_INITIALIZER;
```

```
int pthread_cond_signal(pthread_cond_t *cond);  
int pthread_broadcast(pthread_cond_t *cond);  
int pthread_cond_wait(pthread_cond_t *cond,  
                      pthread_mutex_t *mutex);
```

Bounded-Buffer (two CVs)

- There is a finite-sized buffer that producer threads want to add items to ... and consumer threads want to remove items from ... repeatedly
- **Two** kinds of synchronization needed:
 - **Me**—to protect integrity of the buffer
 - **Correctness**—producer must block if buffer is full and consumer must block if buffer is empty...

Example

```
pthread_mutex_t ring_access =  
    PTHREAD_MUTEX_INITIALIZER;
```

// consumer: wait for content

```
pthread_cond_t some_content =  
    PTHREAD_COND_INITIALIZER;
```

// producer: wait for a free slot

```
pthread_cond_t free_slot =  
    PTHREAD_COND_INITIALIZER;
```

Two CVs, ONE lock!

```
pthread_mutex_t ring_access =  
    PTHREAD_MUTEX_INITIALIZER;
```

// consumer: wait for content

```
pthread_cond_t some_content =  
    PTHREAD_COND_INITIALIZER;
```

// producer: wait for a free slot

```
pthread_cond_t free_slot =  
    PTHREAD_COND_INITIALIZER;
```

Example (cont'd)

```
item_t remove_item (buffer *b){
    item_t st;
    pthread_mutex_lock (&mtx);
    while (b->next_slot_to_retrieve ==
           b->next_slot_to_store)
        pthread_cond_wait(&free_slot, &mtx);

    st = b->items [b->next_slot_to_retrieve];
    b->next_slot_to_retrieve++;
    // adjust next_slot_store if needed
    pthread_cond_signal(&some_content);
    pthread_mutex_unlock (&mtx);
    return st;
}
```

Example (cont'd)

```
void insert_item (buffer *b, item_t st) {  
    pthread_mutex_lock (&mtx);  
    while (b->next_slot_to_store == MAX)  
        pthread_cond_wait (&some_content, &mtx);  
        ...  
    // insert_item  
    pthread_cond_signal (&free_slot);  
    pthread_mutex_unlock (&mtx);  
}
```

BB in action

```
void insert_item (buffer *b, item_t st) {  
    pthread_mutex_lock (&mtx);  
    while (b->next_slot_to_store == MAX)  
        pthread_cond_wait (&some_content, &mtx);  
  
    // insert_item  
    pthread_cond_signal(&free_slot);  
    pthread_mutex_unlock (&mtx);  
}
```

```
item_t remove_item (buffer *b){  
    item_t st;  
    pthread_mutex_lock (&mtx);  
    while (b->next_slot_to_retrieve ==  
    ...  
        b->next_slot_to_store)  
        pthread_cond_wait(&free_slot, &mtx);  
    st = b->items [b->next_slot_to_retrieve];  
    b->next_slot_to_retrieve++;  
    // adjust next_slot_store if needed  
    pthread_cond_signal(&some_content);  
    pthread_mutex_unlock (&mtx);  
    return st;  
}
```


Efficiency

- Taking turns

```
pthread_mutex_t L = PTHREAD_MUTEX_INITIALIZER;;
pthread_cond_t CV = PTHREAD_COND_INITIALIZER;;
int turn = 0;
void* ring (int my_id) {
    while (1) {
        pthread_mutex_lock (&L);
        if (turn == my_id) {
            printf ("%d,", my_id);
            turn = (turn + 1) % N;
            pthread_cond_broadcast (&CV);
        }
        else pthread_cond_wait (&CV, &L);
        pthread_mutex_unlock (&L);
    }
}
```

Locking

- Serializes access
- Two optimizations
 - Granularity
 - Reader-writer

Granularity

- Expedia.com
 - lock (all_hotels_airlines_rental_car_DB)
 - lock (all_hotels)
 - lock (hilton_minneapolis)

Reader Writer

- Now suppose 1000 **just looking** at Hilton minneapolis
 - price
 - look of the rooms

Reader-writer

- Implement using CVs

```
num_readers=0; // 0 free, -1 writer has it
Lock L;
Condition CV;
```

```
acquire_rlock () {
    lock (&L);
    if (num_readers < 0)
        wait (&CV, &L);
    else
        num_readers++;
    unlock (&L);
}
```

```
release() {
    lock (&L);
    if (num_readers == -1)
        num_readers = 0;
    else
        num_readers--;
    if (num_readers == 0)
        broadcast (&CV, &L);
    unlock (&L);
}
```

Write `acquire_wlock` on your own.

```
num_readers=0; // 0 free, -1 writer has it
Lock L;
Condition CV;
```

```
acquire_rlock () {
    lock (&L);
    if (num_readers < 0)
        wait (&CV, &L);
    else
        num_readers++;
    unlock (&L);
}
```

```
release() {
    lock (&L);
    if (num_readers == -1)
        num_readers = 0;
    else
        num_readers--;
    if (num_readers == 0)
        broadcast (&CV, &L);
    unlock (&L);
}
```

Problems with the implementation?

In POSIX

```
#include <pthread.h>

pthread_rwlock_t  L;

pthread_rwlock_rdlock (pthread_rwlock_t *);
pthread_rwlock_wrlock (pthread_rwlock_t *);
pthread_rwlock_unlock (pthread_rwlock_t *);
```


Semaphore

- Synchronization tool does not require busy waiting
- Semaphore operations:

`create_sem`: creates semaphore

`init_sem(ivalue)`: set *value* of semaphore to `ivalue`

`P()`: atomic and indivisible (down/wait)

`V()`: atomic and indivisible (up/post)

`P`: if value is 0 block, otherwise decrement value

`V`: increments *value*, release if anyone blocked

- Internally semaphore structure maintains
 - *Value* // semaphore value, always ≥ 0
 - *Queue* // list of threads waiting in `P()` for the value to be >0

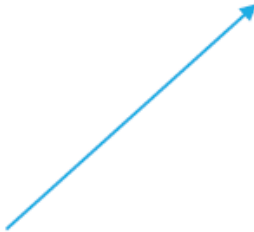
Example

- Counting semaphore example:
 - suppose there are N free resources, n threads ($n > N$)

```
semaphore S: //create
pthread_t t[MaxT]
init_sem (&S,N);
//create N threads
for (i=1;i<n,i++)
    t[i]=pthread_create(...);
```

```
void *fn (...) {
    ...
    P (&S);

    //got resource!
    //do something with it
    V (&S); ...}
```



- Like a lock, it has state!
 - $N=1 \Rightarrow$ binary semaphore \Rightarrow mutual exclusion

Posix Semaphores

```
#include <semaphore.h>

int sem_wait(sem_t* sem); //like P or down
int sem_post(sem_t* sem); // like V or up

//pshared=0 => only threads of process can access

int sem_init(sem_t* sem, int pshared,
             unsigned value);

sem_t sem; //this is akin to create
```

BB with semaphores

//BB of size N with semaphores

```
sem_t consumer_slots, producer_slots;
```

```
sem_init (&consumer_slots, 0, 0);
```

```
sem_init (&producer_slots, 0, N);
```

BB: Posix semaphore (cont'd)

```
void buffer_insert(item_t item) {  
    sem_wait(&producers_slots); //this is like a P()  
    pthread_mutex_lock(&ring_access);  
    ...  
    pthread_mutex_unlock(&ring_access);  
    sem_post(&consumer_slots); //this is like a V()  
}
```

Semaphores using Condition Variables

```
condition CV;  
lock mutex;  
void P() {  
    acquire(&mutex);  
    while (value == 0)  
        wait(&CV, &mutex);  
    value = value - 1;  
    release(&mutex);  
}
```

Deadlock and Synchronization

Dining philosopher

```
while (1) {  
    get_forks();  
    eat();  
    put_forks();  
}
```

