

CSci 4061

Introduction to Operating Systems


Network Systems Programming

Chapter 4.4, 18 (mostly 18.7 onwards) R&R

Goal:

- Enable communication between processes on different machines ... (finally)

Network Layers



advantage of layers?

- Physical Layer
 - media: cat5, radio, transmits bits
 - electrical, data, rates, connectors
- Link layer
 - 802.3 (Ethernet), 802.11 (wifi) transmits error-free
 - framing, CRC codes, ...
- Network layer
 - Routing (path from machine A to machine B)
 - Send packets; packets contain addresses
- Transport layer: systems-programming interface

Protocols

- Transport protocols

- TCP

- Out of order packet assembly
 - Re-transmits packets lost or corrupted
 - Connection-oriented (expensive)
 - Reliable

phone call

- UDP

- Connectionless
 - Unreliable

mailing a letter


- Application protocols

- Applications

- Data formats
 - Meaning of data
 - ...ftp,http,etc

Some network app/protocols

- E-mail
- Web
- Instant messaging
- Remote login
- P2P file sharing
- Multi-user network games
- Streaming stored video clips
- Voice over IP
- Real-time video conferencing



App-layer protocol defines

- Types of messages exchanged
 - e.g. request, response
- Message syntax
 - What fields in messages & how fields are delineated
- Message semantics
 - Meaning of information in fields
- Rules for when and how processes send & respond to messages

HTTP Example

Request:

GET <http://www.cnn.com/> HTTP/1.0

Accept: text/html If-Modified-Since:

Saturday, 15-January-2000 14:37:11

GMT

User-Agent: Mozilla/4.0

Response:

HTTP/1.0 200 OK

Date: Sat, 15 Jan 2000 14:37:12 GMT

Server: Microsoft-IIS/2.0

Content-Type: text/HTML

Content-Length: 1245

Last-Modified: Fri, 14 Jan 2000
08:25:13 GMT

This is the webpage contents

Addresses

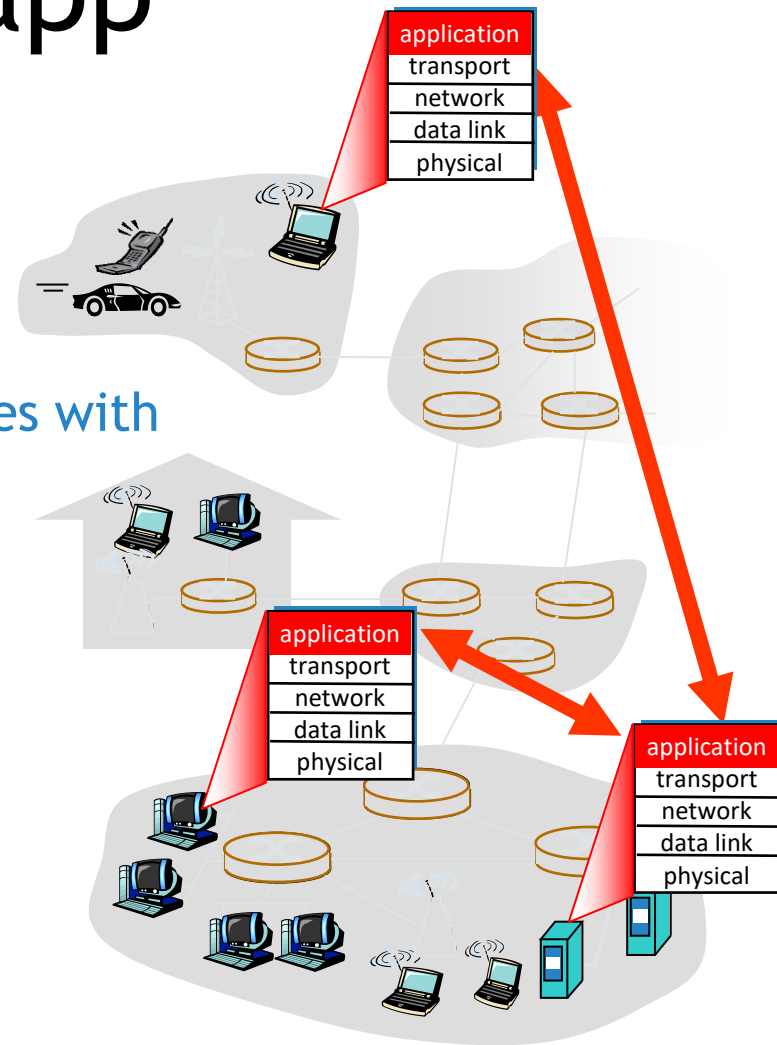


- Local IPC—easy? pid, mailbox, etc
- Network communication: IP address
 - Machine + network interface
 - IPv4 a.b.c.d: 2^{32} , $\sim 4 \times 10^9$ addresses
- Network layer routes packet to destination address (a.b.c.d)
- Symbolic names can be assigned to addresses to identify domains and end hosts
 - *caesar.cs.umn.edu* resolves to a.b.c.d by domain name service (DNS)

Creating a network app

write programs that


- run on (different) *end systems*
- communicate over network
- e.g., web server software communicates with browser software



Application architectures

- Client-server
- Peer-to-peer (P2P)

Client-server architecture



server:

- always-on host
- permanent IP address
- server farms for scaling


clients:

- communicate with server (know its IP address)
- may be intermittently connected
- may have dynamic IP addresses
- why does a client have an IP address?
- do not communicate directly with each other

Pure P2P architecture

- *no* always-on server
- arbitrary end systems “directly” communicate
- peers are intermittently connected and change IP addresses could change
- Highly scalable but difficult to manage

--why?



Systems programming:

Processes communicating

Client process: process that initiates communication

Server process: process that waits to be contacted


- Processes in different hosts communicate by exchanging messages, packets are message fragments

Abstraction:

- Socket - abstraction of a communication endpoint
- <picture>

Sockets

- Process sends/receives messages to/from its **socket**
- socket analogous to door
 - sending process shoves message out the door
 - sending process relies on transport infrastructure on other side of door which brings message to socket at receiving process
- API:
 - (1) choice of transport protocol
 - (2) ability to fix a few parameters
- TCP or UDP sockets



Addressing processes

- To receive messages, process must have an address or *identifier*
- Host device has unique 32-bit IP address (well, could have >1 , but ignore for now)
- Q: Does IP address of host on which process runs suffice for identifying the process?

Addressing processes

- To receive messages, process must have address or *identifier*
 - *Identifier includes both IP address and port numbers associated with process on host.*
- Host device has unique 32-bit IP address
 - Example port numbers:
 - HTTP server: 80
 - Mail server: 25
 - To send HTTP message to caesar.cs.umn.edu web server:
 - IP address: 128.119.245.12
 - Port number: 80

Internet transport protocols services

TCP service:

- ***Connection-oriented***: setup required between client and server processes
- ***Reliable transport*** between sending and receiving process
- ***Does not provide***: timing, minimum throughput guarantees, security

UDP service:

- **Connectionless**: does not provide: **connection** setup, reliability, throughput guarantee, security
- **unreliable** data transfer between sending and receiving process

Message oriented

Q: why bother? Why is there a UDP?

UDP

- Packet is sent as soon as the host sends it!
- TCP adjusts rate of transmission based on congestion control signals