

CSci 4061

Introduction to Operating Systems

Network Systems Programming

Network Programming

Socket headers

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <netdb.h>
```

Creating a socket

```
int socket (int domain, int type, int protocol);
```

- Creates a communication end-point, returns a `fd`
- Domain is the communication domain:
 - `PF_UNIX` is for local Unix communication
 - `PF_INET/INET6` is for IPv4/v6 Internet communication

Creating a socket (cont'd)

- Type specifies the communication style
- TCP
 - `SOCK_STREAM`
 - reliable stream of bytes
- UDP
 - `SOCK_DGRAM`
 - unreliable message delivery

```
fd = socket (PF_INET, SOCK_STREAM, 0);
```

Using sockets: client-server

- Server - 4 steps
 - Creates an open socket with `socket`
 - Bind the socket to a network address (IP, port) with `bind`
 - Set up a queue for incoming connection requests with `listen`
 - `accept` a connection

Using sockets (cont'd)

- Accept connection requests with `accept`
 - Creates a **new** socket connected to the client
 - Returns new socket identifier as an `fd`
- Read and write to/from the new socket `fd` with our old friends `read/write`
- Close the new socket `fd` with `close()` when you are done
- **<picture>**

Socket addressing

- `bind()` takes a socket address structure with an address format that depends on domain

```
int bind (int sockfd, struct sockaddr *my_addr,  
         socklen_t addrlen);
```

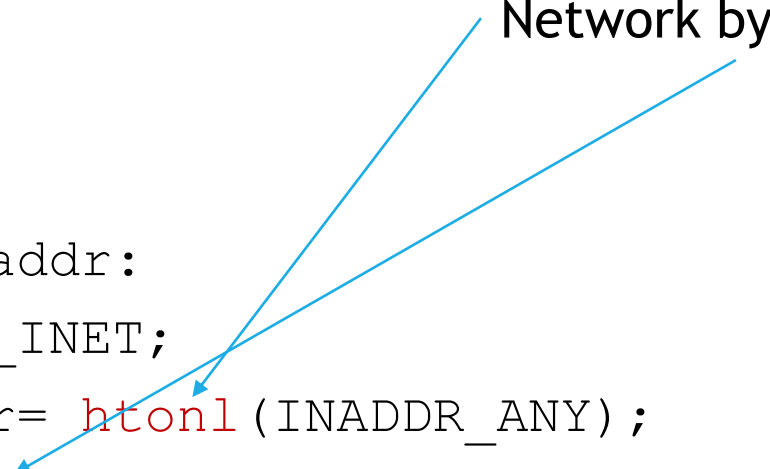
len? different sizes for different types (IPv4, IPv6, ...)

```
struct sockaddr{  
    sa_family_t sin_family;    // AF_INET, same as PF_INET  
    u_int16_t sin_port;        // port: network byte order  
    struct in_addr sin_addr;    // IP address: network byte order  
}
```


Using bind

```
int fd;  
int port = 6666;  
fd = socket (...);  
struct sockaddr_in addr;  
addr.sin_family= AF_INET;  
addr.sin_addr.s_addr= htonl (INADDR_ANY);  
addr.sin_port= htons (port); //server picks the port  
bind (fd, (struct sockaddr*)&addr, sizeof(addr));
```

Network byte order



INADDR_ANY:

OS picks **local IP address**—if multi-homed, then may want to hand-pick address.

Bind error: this port is already taken on this machine

Avoiding Port Collision

- If you are testing your server... `bind` ... then crash ...
- OS holds port for a while and you may get a `bind` error next time you run the server

```
int enable = 1;
```

```
setsockopt(fd, SOL_SOCKET, SO_REUSEADDR,  
           (char*)&enable, sizeof(int));
```

- Gives back ports to OS: call before `bind` so next `bind` call does not fail

Choosing Ports

- You can pick # >1024 (up to 65536)
- `ss` will list active ports

Listen and accept


`listen()` sets up a queue for incoming connection requests

`accept()` gives you a channel to the requester

```
#include <sys/types.h>
#include <sys/socket.h>
int listen (int fd, int backlog);
int accept (int fd, struct sockaddr* addr;
            socklen_t* addrlen);

// create a socket ... enable/reuse ... bind
listen (fd, 5); //queue up 5 pending requests
int accept (fd, (struct sockaddr*)&client_addr
            &addr_len);

//can read or write to new_fd so can the other side
```



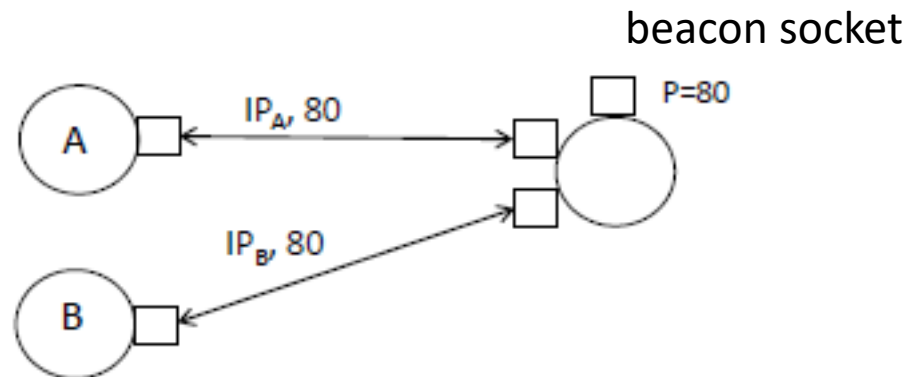
Accept

`accept` returns a `fd` for a new socket

`accept` blocks until a request arrives

Server Styles

- Single thread, one request at a time ...
- Thread/request
 - Each new request is handled by a new thread
 - Could also have a thread pool



Server Styles (cont'd)

- Multiplexed
 - Each thread does a portion of request handling (i.e. lab #3)
- Server cannot selectively `accept` clients!
- Cannot say, `accept` from IP 128.118.44.3
 - Can, however, block/pass it through with a firewall

Now to the client ...

- First, the client needs to know the IP and port of the server
- It can use DNS to get the IP based on the symbolic name
 - [cnn.com](#), [caesar.cs.umn.edu](#), ...

Host Lookup

```
#include <netdb.h>
```


```
struct hostent* gethostbyname(const char *name);
```

“caesar.cs.umn.edu”



hostent **contains** sockaddr

150.91.115.42



Client Scenario

- Create socket with `socket`
- Resolve IP address, port of server
- No need to `bind`, OS does it automatically (when you `connect`) and finds a free port
- Connect to the server with `connect`
- Read and write to/from the socket with
 - `read/write` or other socket I/O calls
 - close the socket with `close`

Internet Addresses

- Internet addresses
 - IP + port
- Must be understandable by the network layer (routers) and end-host
- Network byte order
- Functions to convert components from host to network formats and *vice-versa*

```
#include <inet.h>
```

```
uint32_t htonl(uint32_t hostlong); // long for IP
```

```
uint32_t ntohl(uint32_t netlong);
```

- short int conversion for ports